




Every AI platform has advanced capabilities. None helps you decide how to use them.


Actors



Most Claude users are using just 10% of what it can do prompting it like a search engine



Current Flow



The gap exists at two points: before execution (no guidance) and after output (no transparency)

The 90% use AI like a search engine with better grammar. The 10% gave it access to their actual workflow.



Observation
Every platform hands you powerful tools then leaves you alone to figure out when and how to use them. The more powerful the stack, the more expensive that silence becomes and nobody has more to lose from it than Claude.

Platform	Capabilities beyond Prompting	Strength	Guides Before Executing
Claude	Skills, Agents, Workflows	Most structured and layered capability stack of all four	No, same empty box as everyone else, despite having the deepest capability stack
Chat GPT	GPTs, memory, canvas, plugins	Largest user base, most familiar interface	No, you have to manually search the GPT store yourself and hope you find the right one
Gemini	Gems, Google Workspace extensions	Deeply connected to Gmail, Docs, Drive	No, Gems sit in a sidebar, invisible unless you already know to look
Microsoft Co-Pilot	Copilot agents, plugins, Office suite integration	Context-aware within Word, Excel, Teams	No, Only inside Office apps — step outside those and you're on your own

Observation

Every platform gives you powerful tools. None of them tell you when to use them. The tools exist, the users exist but the moment where one meets the other never comes.

What Data Says

Most people using Claude are still using 10% of what it can do. They prompt it like a search engine.



9 in 10 Would Slow Down For Guidance.

Signals from the Survey

- **80%** → felt unsure if better approach existed
- **90%** → will trade speed for guidance
- **73%** → wasted time from the wrong approach
- **60%** → rewrote the prompt and didn't consider the method change

Existing Evidence

- **"It was okay but needed changes"** 70% said this was their first reaction
- **"Rewrote the prompt"** 60% did this when output failed.
- **"Knowing a different method existed"** 33% said this was missing in that moment.

Secondary Research

- **77%** of IT leaders found high value AI features on their own tools without knowing they existed
- **74%** of workers abandon AI tools mid-task, showing demand for structured guidance over unassisted trial and error
- **63%** of users trust AI more when it explains its reasoning before giving an answer



Hypothesis

- Users blame their prompt, not their approach; when output fails, they rewrite the question, never question the method.
- Awareness doesn't mean action; most knew structured methods existed, almost none reached for them when it mattered.
- "Good enough" kills curiosity; a 70% output is just enough to stop people from asking if better was possible.
- Guidance beats speed, always; 9 in 10 will slow down if something tells them it's worth it.

TAM - 400M+ people use AI chat tools globally
SAM - 10M Claude active users, 77% in the aware but stuck segment = **7.7M addressable users**

✶ They Know It Exists. They Never Use It.

User Segment

Aware but stuck AI users doing real, complex work.

People know structured methods exist. But in the moment of doing the task, none of that knowledge shows up. They default to typing, get a 70% output, rewrite the prompt, and move on.

Why this User Segment?

Because here's where the solution actually works.

Someone who doesn't know Skills exist needs education first. Someone who uses structured methods regularly doesn't need a nudge. This segment sits exactly in the middle: 46% tried once or twice, 31% use regularly hitting the same wall. That's 77% of respondents in this zone.

User Persona 01

Anuvrat 22

About -

An early professional who uses AI almost daily for tasks that require real thinking like project reports, code debugging, research summaries, automating parts of his workflow. He takes his outputs seriously.



Goals

- Get a good output without overthinking the tool
- Feel confident what they produced was actually good

Needs

- To be told right now a better path exists
- Proof it'll be better, not just different

Blocker

- Nothing tells them when a better approach would help
- "Good enough" output stops them from questioning

Pain Points

- Spends more time fixing outputs than he should
- Tried Skills once, setup felt like extra work

User Persona 02

Swati 42

About -

A content writer who uses Claude every single day for writing, research, summarizing, and drafting. She is not here to explore the tool, she is here to get work done.



Goals

- Get clean, usable content output without spending time figuring out the tool
- Feel confident the draft she got was actually the best Claude could give her

Needs

- To know in the moment not after, that a better approach exists
- Reassurance that switching won't break her flow or waste her time

Blocker

- always mid-task when she would need to make a capability decision
- Every time she has tried something beyond prompting it felt like a detour.

Pain Points

- Spends time re-prompting outputs that a Skill would have handled better
- Has no way of knowing whether what she got was the best possible

✶ No Layer Between Intent and Output.

What is the True Problem?

There is **no decision layer** between a user's intent and Claude's execution. Claude **decides silently** which capability to use. The **user never sees it**, never learns from it, never gets better at using the tool. The **system was built for speed**, and it removed user agency in the process.



Who are the customer facing problem?

Aware but stuck AI users doing real, complex work. **77%** of the respondents in the zone of knowing structured methods exist but not reaching for them when it matters.

How do we know it is a real problem?

From the survey:

→ **80%** felt **unsure** a better approach existed in the moment

From external survey:

"Most people barely scratch the surface of what Claude can do. I used to treat Claude like a glorified search engine quick questions, one-line prompts."

What is the value generated by solving this problem?

For Users

- They stop wasting time rewriting prompts and get to the right approach the first time
- They feel confident the output they submitted was actually the best possible not just "good enough"

For Business

- For Anthropic, Skills were built to drive Pro upgrades. Without discovery, that investment generates zero return.
- Users who use advanced features retain better and upgrade plans.



Why should we solve this problem now?

- Skills launched October 2025 and are now active across all plans, the feature is live, the discovery layer doesn't exist yet.
- Bad habits form fast. Every week users spend defaulting to basic prompting is another week that behaviour gets harder to change.



- Solution • 01

Echo

 - Claude declares which capability it activated and why, in one line
 - Zero friction, no trigger, no decision required from the user
 - Makes the invisible visible in every response, every session
- Solution • 02

GroundWork

 - Claude states its intended approach before executing
 - User confirms, corrects, or proceeds
 - This fixes the root cause; silent decisions become collaborative ones
- Solution • 03

Crux

 - Claude opens with its confident on the task
 - Names the one missing input that would change the output
 - Turns trust from a barrier into an active conversation.

Solution Selected : The GroundWork

Fixes the root cause, not the symptom
→ Every other solution catches the problem mid-flow or after. This one lives exactly where the 5 Whys pointed that was the missing decision layer at input.

Highest trust signal of all five
→ Claude showing its reasoning before executing directly addresses why 80% felt unsure. It explains itself.

Foundation for future expansion
→ GroundWork as the core, then layer:
1. Crux for the trust gap
2. Echo for passive transparency on every response

Solution	Reach	Impact	Confidence	Effort	RICE Score
Echo	10	2	65%	1	13.0
GroundWork	7	3	75%	2	7.9
Crux	8	2	70%	1.5	7.5

Reason

While Echo had the highest RICE score, prioritized GroundWork because it addresses the core issue of visibility just before the execution and not after. Echo catches the problem after output; GroundWork intercepts before the decision is made.

Declare The Plan. Then Execute.

What Exactly is it?

A chat interface where Claude tells you what it's about to do before it does it. For bigger tasks, it also tells you if there's a Skill that could do it better and lets you compare both outputs before you commit to one.

What It is NOT?

- Not a place to build or edit Skills
- Not a prompting guide or tutorial
- Not something you set up or configure, it just works from message one

A decision layer that makes Claude's approach visible before it executes.

What GroundWork Does

- **Approach Declaration Card:** Claude tells you exactly how it plans to tackle your task before it writes a single word
- **Smart Task Classification:** every prompt is automatically read and sorted into simple, medium, or complex and the experience changes accordingly
- **Skill Suggestion with Context:** when a Skill exists for your task, Claude tells you what it is, what it does, and why it might help, right before executing
- **Side by Side Comparison:** one click generates both a standard response and a Skill powered response simultaneously so you can read both and pick one
- **Skill FAQ Chips:** three clickable questions appear in the chat so you can understand what a Skill is before deciding to use it, without leaving the conversation


See It Before You Send It

 Wire Frame


 Prototype

 Data Flow


 User Journey



1. Clean prompt entry for starting any task.



2. Skill decision layer before execution.



3. Side-by-side outputs to show Skill impact.

1. Clean prompt entry Input centered, zero distraction. Three pre-loaded prompts set expectations without any instructions. Subtitle prepares the user before they type so the card never feels like a surprise.
2. Skill decision layer Decision appears before output so the user still has full agency. FAQ chips let them explore without committing. Adjust and Cancel sit at the same level as Skill buttons so backing out feels as easy as proceeding.
3. Side by side comparison Both responses generate at once, no waiting twice. Go with this one sits at the bottom of each column so picking feels like a conclusion not an interruption.

How it works - from prompt to the right approach

Step 1 - Claude reads the task Every prompt is classified before execution.
Simple → straight to answer. **Medium** → Claude declares its plan. **Complex** → Claude declares its plan and suggests a Skill.

Step 2 - User decides Confirm the approach or adjust it in plain language.
For complex tasks, use the Skill, skip it, or see both side by side. One decision. No setup required.

Step 3 - Claude executes and learns Claude runs the chosen path.
A single thumbs up or down at the end feeds back into the classifier. Every session makes the next suggestion smarter.

Skill Adoption Or Nothing.

North Star Metric

Skill Adoption Rate

Users who chose Use Skill or picked Skill in side by side / Total users who saw a Skill suggestion x 100

L1 Metric

- **Approach Card Confirm Rate** → confirmed / saw card x 100
- **Skill Suggestion Click Rate** → clicked any button / saw suggestion x 100
- **Side by Side Pick Rate** → picked Skill response / used I don't know x 100

L2 Metric

- **Adjust Rate** → adjusted approach / saw card x 100
- **FAQ Chip Click Rate** → clicked a chip / saw suggestion x 100
- **Session Return** → returned users / total users x 100

Leading Indicators

- FAQ Chip Click Rate
- Skill Suggestion Rate
- Adjust Rate

early signals that show the prototype is changing behaviour before outcomes are visible

Lagging Indicators

- Skill Adoption Rate
- Session Return Rate
- Side by Side Pick Rate

confirmation signals that show whether the prototype actually moved the needle over time

KPI Tree



Monetisation

Now → Retention

Skills are **already in Pro and Team plans** but not a lot of people uses them. This prototype makes people actually use them. That alone **makes the paid plan feel worth it** and stops people from leaving.

Mid term → Upgrades

When a free user sees the side by side comparison and **notices the Skill response** is clearly better, that is the moment they consider upgrading. **No popup, no ad.** The product sells itself.

Long term → Skill Tiers

Once we know which Skills people love using the most, those become premium. **Usage data** from the **feedback flywheel** tells us exactly which ones to charge for.

7.7M addressable users, 1% conversion at Rs 299 per month equals Rs 23 crore ARR potential.



Every Risk Has A Built-In Fix.

Risk	Mitigation Strategy
Users ignore the approach card and just click Confirm without reading	Make the card short and scannable, one line only, low effort to read
Users find the extra step annoying and drop off	Simple tasks get no card at all, friction only appears when it adds value
Skill suggestions feel irrelevant or random	Classifier improves over time using the feedback flywheel data
Side by side comparison confuses users instead of helping	FAQ chips give context before the choice, user is never left guessing
Low feedback submission rate, flywheel never starts	Thumbs up and down are one click, no text, no friction, lowest possible barrier

What's Next?

v1.1 - Rewind If a **user skips the Skill** suggestion and the **output disappoints** them, Claude catches that **signal and offers the Skill** as a second chance. No friction the first time, safety net the second time.

v2 - Personalised Routing The **classifier learns your patterns** over time. If you always confirm the approach and always pick the Skill response, **Claude stops asking and just does it**. The **decision layer becomes invisible** because it already knows you.

v3 - Skill Genesis Claude **notices when you keep doing the same type of task** and asks **if it should save your preferred approach as a Skill**. One click and it is yours. **No setup, no prompt engineering**. Your **behaviour becomes the Skill**.

Experiment Plan

Show the Approach Card to half the users, hide it from the other half. Run for 2 weeks across 1000 users. If the group that saw the card adopted Skills 15% more, roll out to everyone. If the difference is under 5%, redesign the card.